

# CSCI 2033 Assignment 3:

## QR Decomposition

Posted: Saturday, November 18

Updated: Wednesday, November 22

Due date: Thursday, November 30, 11:55 pm

In this assignment, you will implement a MATLAB function to decompose a matrix  $\mathbf{A}$  into the product of two matrices  $\mathbf{A} = \mathbf{QR}$ , where  $\mathbf{Q}$  has orthonormal columns and  $\mathbf{R}$  is upper triangular. You will then use your decomposition to solve a shape fitting problem.

You will need the files `back_sub.m` for Part 4 and `generate_data.m` and `visualize.m` for Part 5. These can be found in the zip file provided on the Moodle assignment page.

## 1 Submission Guidelines

You will submit a zip file that contains the following `.m` files on Moodle:

- `ortho_decomp.m`
- `my_qr.m`
- `least_squares.m`
- `my_pack.m`
- `my_unpack.m`
- `design_matrix.m`
- `affine_fit.m`

as well as any helper functions that are needed by your implementation.

As with previous assignments, please note the following rules:

1. Each file must be named as specified above.
2. Each function must return the specified value(s).
3. You may use MATLAB's array manipulation syntax, e.g. `size(A)`, `A(i,:)`, and `[A,B]`, and basic operations like addition, multiplication, and transpose of matrices and vectors. High-level linear algebra functions such as `inv`, `qr`, and `A\b` are *not* allowed except where specified. Please contact the instructor with further questions.

4. This is an individual assignment, and no collaboration is allowed. Any in-person or online discussion should stop before you start discussing or designing a solution.

## 2 Orthogonal decomposition

Suppose you have an orthonormal basis  $\{\mathbf{u}_1, \dots, \mathbf{u}_n\}$  for a subspace  $H \subseteq \mathbb{R}^m$ . Any vector  $\mathbf{v} \in \mathbb{R}^m$  can be decomposed into the sum of two orthogonal vectors,  $\hat{\mathbf{v}} \in H$  and  $\hat{\mathbf{v}}^\perp \in H^\perp$ . Furthermore, since  $\hat{\mathbf{v}} \in H$ , it can be expressed as  $\hat{\mathbf{v}} = c_1\mathbf{u}_1 + \dots + c_n\mathbf{u}_n$  for some weights  $c_1, \dots, c_n$ . Implement a function to perform this decomposition.

**Specification:**

**function** `[c, v_perp] = ortho_decomp(U, v)`

**Input:** an  $m \times n$  matrix  $\mathbf{U} = [\mathbf{u}_1 \ \dots \ \mathbf{u}_n]$  with orthonormal columns, and a vector  $\mathbf{v}$ .

**Output:**

**c:** a vector  $\mathbf{c} = \begin{bmatrix} c_1 \\ \vdots \\ c_n \end{bmatrix}$  such that  $\mathbf{v} = c_1\mathbf{u}_1 + \dots + c_n\mathbf{u}_n + \hat{\mathbf{v}}^\perp$

**v\_perp:** the residual vector  $\mathbf{v}^\perp$  such that  $\mathbf{u}_i \cdot \hat{\mathbf{v}}^\perp = 0$  for all  $i$ .

**Implementation notes:**

Here  $\hat{\mathbf{v}}$  is simply the orthogonal projection of  $\mathbf{v}$  onto the subspace  $H$ . Since we have an orthogonal basis of  $H$ , this should be easy to compute. In fact, since the basis is *orthonormal*, it is possible to compute the projection in just one line of code without any `for` loops. Hint: What are the entries of  $\mathbf{U}^T \mathbf{v}$ ?

**Test cases:**

$$\begin{aligned}
 1. \quad \mathbf{U} &= \begin{bmatrix} 1 & 0 \\ 0 & 0 \\ 0 & 1 \\ 0 & 0 \end{bmatrix}, \mathbf{v} = \begin{bmatrix} 2 \\ 3 \\ 4 \\ 5 \end{bmatrix} \quad \rightarrow \quad \mathbf{c} = \begin{bmatrix} 2 \\ 4 \end{bmatrix}, \hat{\mathbf{v}}^\perp = \begin{bmatrix} 0 \\ 3 \\ 0 \\ 5 \end{bmatrix} \\
 2. \quad \mathbf{U} &= \begin{bmatrix} 0.6 \\ 0.8 \\ 0 \end{bmatrix}, \mathbf{v} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} \quad \rightarrow \quad \mathbf{c} = [1.4], \hat{\mathbf{v}}^\perp = \begin{bmatrix} 0.16 \\ -0.12 \\ 1 \end{bmatrix}
 \end{aligned}$$

3. Generate a random orthonormal set of 5 vectors in  $\mathbb{R}^{10}$  using the command `[U, ~] = qr(randn(10,5),0)`, and generate another random integer vector `v = randi([-2,2], 10,1)`. Compute your orthogonal decomposition, `[c, v_perp] = ortho_decomp(U, v)`. Verify that `U*c + v_perp` is nearly the same as `v`, and `U'*v_perp` is nearly zero (both to within  $10^{-12}$ ).

### 3 QR decomposition

Suppose you have an  $m \times n$  matrix  $\mathbf{A}$  which is “tall and thin”, i.e. with  $m > n$ , and the columns of  $\mathbf{A}$  are linearly independent. The Gram-Schmidt process corresponds to a factorization  $\mathbf{A} = \mathbf{QR}$ , where  $\mathbf{Q}$  is an  $m \times n$  matrix with orthonormal columns, and  $\mathbf{R}$  is an  $n \times n$  upper triangular matrix.<sup>1</sup>

#### Specification:

**function** `[Q, R] = my_qr(A)`

**Input:** an  $m \times n$  matrix  $\mathbf{A}$  with  $m > n$ .

**Output:** an  $m \times n$  matrix  $\mathbf{Q}$  with orthonormal columns, and an  $n \times n$  upper triangular matrix  $\mathbf{R}$ , such that  $\mathbf{A} = \mathbf{QR}$ .

#### Implementation notes:

I recommend starting your implementation by first having your function compute  $\mathbf{Q}$  correctly. After that, you can modify your code to also compute  $\mathbf{R}$ .

The columns of  $\mathbf{Q}$  are essentially obtained by performing the Gram-Schmidt process with normalization on the columns of  $\mathbf{A}$ :

$$\begin{aligned}
 \mathbf{q}_1 &= \mathbf{a}_1 / \|\mathbf{a}_1\|, \\
 \hat{\mathbf{a}}_2^\perp &= \mathbf{a}_2 - \text{proj}_{H_1} \mathbf{a}_2 & \text{where } H_1 = \text{span}\{\mathbf{q}_1\}, \\
 \mathbf{q}_2 &= \hat{\mathbf{a}}_2^\perp / \|\hat{\mathbf{a}}_2^\perp\|, \\
 \hat{\mathbf{a}}_3^\perp &= \mathbf{a}_3 - \text{proj}_{H_2} \mathbf{a}_3 & \text{where } H_2 = \text{span}\{\mathbf{q}_1, \mathbf{q}_2\}, \\
 \mathbf{q}_3 &= \hat{\mathbf{a}}_3^\perp / \|\hat{\mathbf{a}}_3^\perp\|, \\
 &\vdots
 \end{aligned}$$

You can carry this out using your orthogonal decomposition function. For each column  $i = 1, \dots, n$ , call `ortho_decomp` with an appropriate set of inputs, and

---

<sup>1</sup>Technically, what we are computing here is known as the “thin” or “reduced” QR decomposition. In the full QR decomposition,  $\mathbf{Q}$  is square and orthogonal, and  $\mathbf{R}$  is an  $m \times n$  upper triangular matrix whose lower  $(m - n)$  rows are all zero. In practice, it is also not computed with the Gram-Schmidt process, but with other methods that are more numerically stable.

use the resulting  $\hat{\mathbf{v}}^\perp$  to fill in the  $i$ th column of  $\mathbf{Q}$ . In the end, the columns of  $\mathbf{Q}$  should be an orthonormal set of vectors  $\{\mathbf{q}_1, \dots, \mathbf{q}_n\}$  which span  $\text{Col } \mathbf{A}$ .

Once you can compute  $\mathbf{Q}$  correctly, modify your function to compute  $\mathbf{R}$  as well. In principle, since  $\mathbf{Q}^T \mathbf{Q} = \mathbf{I}$ , you could obtain  $\mathbf{R}$  simply via  $\mathbf{Q}^T \mathbf{A} = \mathbf{Q}^T \mathbf{Q} \mathbf{R} = \mathbf{R}$ . However, this is more work than necessary, because you can actually fill in the entries of  $\mathbf{R}$  as you compute each column of  $\mathbf{Q}$ . When you compute the  $i$ th column of  $\mathbf{Q}$ , the `ortho_decomp` call gives you both  $\mathbf{c}$  and  $\hat{\mathbf{v}}^\perp$ ; can you use these to fill in the  $i$  nonzero entries in the  $i$ th column of  $\mathbf{R}$ ? Hint: Consider the case

when  $\mathbf{A}$  is already an upper triangular matrix, say  $\mathbf{A} = \begin{bmatrix} 1 & 2 & 3 \\ 0 & 4 & 5 \\ 0 & 0 & 6 \end{bmatrix}$ , and work out by hand what  $\mathbf{c}$  and  $\hat{\mathbf{v}}^\perp$  will be for each column.

**Test cases:**

$$1. \mathbf{A} = \begin{bmatrix} 0 & 3 & 0 \\ 0 & 0 & -4 \\ -2 & 0 & 0 \end{bmatrix} \rightarrow \mathbf{Q} = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & -1 \\ -1 & 0 & 0 \end{bmatrix}, \mathbf{R} = \begin{bmatrix} 2 & 0 & 0 \\ 0 & 3 & 0 \\ 0 & 0 & 4 \end{bmatrix}$$

$$2. \mathbf{A} = \begin{bmatrix} 1 & 1 & 0 \\ 1 & 0 & 0 \\ 1 & 1 & 1 \\ 1 & 0 & 1 \end{bmatrix} \rightarrow \mathbf{Q} = \begin{bmatrix} 1/2 & 1/2 & -1/2 \\ 1/2 & -1/2 & -1/2 \\ 1/2 & 1/2 & 1/2 \\ 1/2 & -1/2 & 1/2 \end{bmatrix}, \mathbf{R} = \begin{bmatrix} 2 & 1 & 1 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

3. Generate a random  $10 \times 5$  integer matrix  $\mathbf{A} = \text{randi}([-2, 2], 10, 5)$  and compute your QR decomposition,  $[\mathbf{Q}, \mathbf{R}] = \text{my\_qr}(\mathbf{A})$ . Verify that `istriu(R)` is true,  $\mathbf{Q}' * \mathbf{Q}$  is nearly the identity matrix, and  $\mathbf{Q} * \mathbf{R}$  is nearly the same as  $\mathbf{A}$  (both to within  $10^{-12}$ ).

## 4 Least-squares problems

Use the QR decomposition to solve the least-squares problem  $\mathbf{Ax} \approx \mathbf{b}$ .

### Specification:

`function x = least_squares(A, b)`

**Input:** an  $m \times n$  matrix  $\mathbf{A}$  and a vector  $\mathbf{b} \in \mathbb{R}^m$ .

**Output:** a vector  $\mathbf{x} \in \mathbb{R}^n$  which is the least-squares solution of the over-determined system  $\mathbf{Ax} \approx \mathbf{b}$ , i.e. such that  $\mathbf{Ax} - \mathbf{b} \in (\text{Col } \mathbf{A})^\perp$ .

### Implementation notes:

Do not solve the normal equations  $\mathbf{A}^T \mathbf{Ax} = \mathbf{A}^T \mathbf{b}$  directly, as this is a very numerically unstable approach. Instead, compute the QR factorization of  $\mathbf{A}$  and use it to solve the least-squares problem with only a matrix-vector multiplication and a back-substitution, as described in the textbook. Use the `back_sub` function provided with this assignment to perform the back-substitution.

If you cannot get your `my_qr` function to work, you may use MATLAB's built-in `qr` function here so you can still attempt Part 5. Call `qr(A,0)` to get a rectangular  $\mathbf{Q}$  and square  $\mathbf{R}$  as desired.

### Test cases:

$$1. \mathbf{A} = \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix}, \mathbf{b} = \begin{bmatrix} 2 \\ 3 \\ 4 \end{bmatrix} \rightarrow \mathbf{x} = [2.5]$$

$$2. \mathbf{A} = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 1 \\ 1 & 2 & 4 \\ 1 & 3 & 9 \\ 1 & 4 & 16 \end{bmatrix}, \mathbf{b} = \begin{bmatrix} 0 \\ 3 \\ 4 \\ 3 \\ 0 \end{bmatrix} \rightarrow \mathbf{x} = \begin{bmatrix} 0 \\ 4 \\ -1 \end{bmatrix}$$

3. Generate a random  $10 \times 5$  integer matrix  $\mathbf{A} = \text{randi}([-2,2], 10, 5)$  and a random integer vector  $\mathbf{b} = \text{randi}([-2,2], 10, 1)$ . Compute your least-squares solution,  $\mathbf{x} = \text{least\_squares}(\mathbf{A}, \mathbf{b})$ , and obtain the residual  $\mathbf{r} = \mathbf{b} - \mathbf{A}*\mathbf{x}$ . Verify that  $\mathbf{A}'*\mathbf{r}$  is nearly zero (to within  $10^{-12}$ ).

## 5 Best-fitting transformations

An *affine* transformation is a combination of a linear transformation and a translation (i.e. displacement) while retaining parallelism, i.e., parallel lines remain parallel after the transformation. For example, a point  $\begin{bmatrix} x \\ y \end{bmatrix}$  on A (orange square) in Figure 1 can be transformed to  $\begin{bmatrix} \bar{x} \\ \bar{y} \end{bmatrix}$  on B (red parallelogram) via a

linear transform  $\mathbf{M} = \begin{bmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \end{bmatrix}$ , i.e.,

$$\begin{bmatrix} \bar{x} \\ \bar{y} \end{bmatrix} = \begin{bmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}.$$

Note that the parallel lines stay parallel. This transform can be combined with a translation, moving the red parallelogram to blue parallelogram (C). This composite transformation can be written as:

$$\begin{bmatrix} \tilde{x} \\ \tilde{y} \end{bmatrix} = \begin{bmatrix} \bar{x} \\ \bar{y} \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix} = \begin{bmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix} = \mathbf{M} \begin{bmatrix} x \\ y \end{bmatrix} + \mathbf{t}, \quad (1)$$

where  $\mathbf{t} = \begin{bmatrix} t_x \\ t_y \end{bmatrix}$  is the translation vector.




Figure 1: Affine transform.

In sum, Equation (1) can be written as

$$\begin{bmatrix} \tilde{x} \\ \tilde{y} \end{bmatrix} = \begin{bmatrix} m_{11}x + m_{12}y + t_x \\ m_{21}x + m_{22}y + t_y \end{bmatrix}. \quad (2)$$

## 5.1 Linear Equation from a Single Correspondence

Your task is given many correspondences<sup>2</sup>  $(x_i, y_i) \leftrightarrow (\tilde{x}_i, \tilde{y}_i)$  where  $i$  is the index for the correspondence, to compute the best affine transform parameters,  $m_{11}, m_{12}, m_{21}, m_{22}, t_x, t_y$  (unknowns): We can rewrite Equation (2) by arranging it with respect to the unknowns for, say, the first correspondence:

$$\underbrace{\begin{bmatrix} \tilde{x}_1 \\ \tilde{y}_1 \end{bmatrix}}_{\tilde{\mathbf{p}}_1} = \underbrace{\begin{bmatrix} x_1 & y_1 & 0 & 0 & 1 & 0 \\ 0 & 0 & x_1 & y_1 & 0 & 1 \end{bmatrix}}_{\mathbf{A}_1} \underbrace{\begin{bmatrix} m_{11} \\ m_{12} \\ m_{21} \\ m_{22} \\ t_x \\ t_y \end{bmatrix}}_{\boldsymbol{\beta}}. \quad (3)$$

or simply,

$$\tilde{\mathbf{p}}_1 = \mathbf{A}_1 \boldsymbol{\beta}. \quad (4)$$

Note that  $\mathbf{A}_1$  depends on the original position of the point,  $\mathbf{p}_1 = \begin{bmatrix} x_1 \\ y_1 \end{bmatrix}$ . The correspondence produces 2 linear equations while the number of unknowns is 6, so at least 3 correspondences are needed to uniquely determine the affine transform parameters.

Implement a function `beta = my_pack(M, t)` to obtain  $\boldsymbol{\beta}$  from  $\mathbf{M} = \begin{bmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \end{bmatrix}$  and  $\mathbf{t} = \begin{bmatrix} t_x \\ t_y \end{bmatrix}$ , and its inverse function `[M, t] = my_unpack(beta)`.

`function beta = my_pack(M, t)`

**Input:** a  $2 \times 2$  matrix  $\mathbf{M}$  and a vector  $\mathbf{t}$ .

**Output:** a vector  $\boldsymbol{\beta} \in \mathbb{R}^6$  containing the entries of  $\mathbf{M}$  and  $\mathbf{t}$ .

`function [M, t] = my_unpack(beta)`

**Input:** a vector  $\boldsymbol{\beta} \in \mathbb{R}^6$ .

**Output:** a  $2 \times 2$  matrix  $\mathbf{M}$  and a vector  $\mathbf{t} \in \mathbb{R}^2$  such that `my_pack(M, t) = beta`.

Finally, construct the matrix  $\mathbf{A}_i$  given  $\mathbf{p}_i$ .

`function Ai = design_matrix(pi)`

**Input:** a vector  $\mathbf{p}_i = \begin{bmatrix} x_i \\ y_i \end{bmatrix} \in \mathbb{R}^2$ .

**Output:** the matrix  $\mathbf{A}_i$  given by Equation (3).

---

<sup>2</sup>A point  $(x, y)$  in A is transformed to a point  $(\tilde{x}, \tilde{y})$  in C. This pair of points forms a *correspondence*.

## 5.2 Solving Linear System from $n$ Correspondences

Equation (3) can be extended to include  $n$  correspondences as follow:

$$\begin{bmatrix} \tilde{x}_1 \\ \tilde{y}_1 \\ \vdots \\ \tilde{x}_n \\ \tilde{y}_n \end{bmatrix} = \begin{bmatrix} x_1 & y_1 & 0 & 0 & 1 & 0 \\ 0 & 0 & x_1 & y_1 & 0 & 1 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ x_n & y_n & 0 & 0 & 1 & 0 \\ 0 & 0 & x_n & y_n & 0 & 1 \end{bmatrix} \begin{bmatrix} m_{11} \\ m_{12} \\ m_{21} \\ m_{22} \\ t_x \\ t_y \end{bmatrix}, \quad (5)$$

or again simply,

$$\begin{bmatrix} \tilde{\mathbf{p}}_1 \\ \vdots \\ \tilde{\mathbf{p}}_n \end{bmatrix} = \begin{bmatrix} \mathbf{A}_1 \\ \vdots \\ \mathbf{A}_n \end{bmatrix} \boldsymbol{\beta}, \quad (6)$$

If all correspondences are noise-free, Equation (6) will be always satisfied. Due to correspondence noise in practice, it cannot be satisfied, and therefore,

$$\begin{bmatrix} \tilde{\mathbf{p}}_1 \\ \vdots \\ \tilde{\mathbf{p}}_n \end{bmatrix} \approx \begin{bmatrix} \mathbf{A}_1 \\ \vdots \\ \mathbf{A}_n \end{bmatrix} \boldsymbol{\beta}, \quad (7)$$

Now we will compute the best affine transform parameters using Equation (7).

**function** [M, t] = affine\_fit(P, P\_tilde)

**Input:**  $2 \times k$  correspondence matrices  $\mathbf{P} = [\mathbf{p}_1 \ \cdots \ \mathbf{p}_k]$  and  $\tilde{\mathbf{P}} = [\tilde{\mathbf{p}}_1 \ \cdots \ \tilde{\mathbf{p}}_k]$  containing the reference points and transformed points as columns.

**Output:** a  $2 \times 2$  matrix  $\mathbf{M}$  and a vector  $\mathbf{t} \in \mathbb{R}^2$  such that the affine transformation  $\mathbf{M}\mathbf{x} + \mathbf{t}$  best matches the given data. To compute this, you will have to set up and solve the least-squares problem in Equation (7) to obtain  $\boldsymbol{\beta}$ , then unpack it to get  $\mathbf{M}$  and  $\mathbf{t}$  out.


**Test cases:**

For the first two tests, pick an arbitrary  $\mathbf{M}$  and  $\mathbf{t}$ , say  $\mathbf{M} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$  and  $\mathbf{t} = \begin{bmatrix} 5 \\ 6 \end{bmatrix}$ .

1. Check that `[M_new, t_new] = my_unpack(my_pack(M, t))` recovers the original values of  $\mathbf{M}$  and  $\mathbf{t}$ .
2. Verify that when  $\mathbf{x} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$ , `design_matrix(x)*my_pack(M,t)` gives back  $\mathbf{t}$ . Similarly,  $\mathbf{x} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$  and  $\mathbf{x} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$  should give you  $\mathbf{m}_1 + \mathbf{t}$  and  $\mathbf{m}_2 + \mathbf{t}$  respectively, where  $\mathbf{m}_1$  and  $\mathbf{m}_2$  are the columns of  $\mathbf{M}$ .



3. Applying `affine_fit` to  $\mathbf{P} = \begin{bmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 \end{bmatrix}$ ,  $\tilde{\mathbf{P}} = \begin{bmatrix} 0.4 & 1 & 1.8 & 1.2 \\ 0.8 & 0 & 0.6 & 1.4 \end{bmatrix}$  should give  $\mathbf{M} = \begin{bmatrix} 0.6 & 0.8 \\ -0.8 & 0.6 \end{bmatrix}$  and  $\mathbf{t} = \begin{bmatrix} 0.4 \\ 0.8 \end{bmatrix}$ : a rotation and a translation.



4. We have provided a function `[P, P_tilde, M, t] = generate_data()` that produces some random test data. It does so by filling random values in  $\mathbf{M}$  and  $\mathbf{t}$ , choosing random points  $\mathbf{p}_i$ , then setting each  $\tilde{\mathbf{p}}_i$  to  $\mathbf{M}\mathbf{p}_i + \mathbf{t}$  plus a small amount of random noise. You can visualize the data by calling the provided function `visualize(P, P_tilde)`.

Call `affine_fit(P, P_tilde)` to obtain your own estimates of  $\mathbf{M}$  and  $\mathbf{t}$ . The result should be close to the “ground truth” transformation returned by `generate_data`, although due to the added noise it will not be exactly identical. Call `visualize(P, P_tilde, M, t)` to see what your fit looks like.

Note: If you want to test on a smaller data set, just use the first few columns of  $\mathbf{P}$  and  $\tilde{\mathbf{P}}$ , for example  $\mathbf{P} = \mathbf{P}(:, 1:10)$  and  $\mathbf{P\_tilde} = \mathbf{P\_tilde}(:, 1:10)$ .